Relational Databases and Datawarehousing – SQL

Subqueries, Insert-Update-Delete-Merge, Views, Common Table Expressions

SUBQUERIES

SUBQUERIES basic form

- Nested subqueries
  - Basic form

    ```
    SELECT
    FROM
    WHERE condition
    ```

    → Contains in its left and/or right hand side statement another SELECT

    - Outer level query = the first SELECT. This is the main question.
    - Inner level query = the SELECT in the WHERE clause (or HAVING clause). This is the subquery:
      - Always executed first
      - Always between ().
      - Subqueries can be nested at > 1 level.
  - A subquery can return one value or a list of values

SUBQUERY that returns a <u>single value</u>

– The result of the query can be used anywhere you can use an expression.

- With all relational operators: =, >, <, <=,>=,<>

- Example:

  – What is the UnitPrice of the most expensive product?

  ```sql
  SELECT MAX(UnitPrice) As MaxPrice FROM Products
  ```

	MaxPrice
1	263,50

  – What is the most expensive product?

  First the table Products is searched to determine the highest salary (= subquery). Then the table Products is searched a second time (= main query) to evaluate each unitprice against the determined maximum.

  ```sql
  SELECT ProductID, ProductName, UnitPrice As MaxPrice
  FROM Products
  WHERE UnitPrice = (SELECT MAX(UnitPrice) FROM Products)
  ```

	ProductID	ProductName	MaxPrice
1	38	Côte de Blaye	263,50

  – Returns the previous query always the same resultset as the following?

  ```sql
  SELECT TOP 1 ProductID, ProductName, UnitPrice As MaxPrice
  FROM Products ORDER BY UnitPrice DESC
  ```

SUBQUERY that returns a <u>single value</u>

- Other examples
  - Give the products that cost more than average

```
SELECT ProductID, ProductName, UnitPrice As MaxPrice
FROM Products
WHERE UnitPrice > (SELECT AVG(UnitPrice) FROM Products)
```

  - Who is the youngest employee from the USA?

```
SELECT LastName, FirstName
FROM Employees
WHERE Country = 'USA'
AND BirthDate = (SELECT MAX(BirthDate) FROM Employees WHERE Country = 'USA')
```

SUBQUERY that returns a <u>single column</u>

- The resulting column can be used as a list
  - Operators IN, NOT IN, ANY, ALL
  - IN operator (=ANY operator)
    - Example: Give all employees that have processed orders

    ```
    SELECT e.EmployeeID, e.FirstName + ' ' + e.LastName As Name
    FROM Employees e
    WHERE e.EmployeeID IN (SELECT DISTINCT EmployeeID FROM Orders)
    ```

    - This can also be accomplished with JOIN

    ```
    SELECT DISTINCT e.EmployeeID, e.FirstName + ' ' + e.LastName As Name
    FROM Employees e JOIN Orders o ON e.EmployeeID = o.EmployeeID
    ```

SUBQUERY that returns a <u>single column</u>

- The resulting column can be used as a list
  - Operators IN, NOT IN, ANY, ALL
  - NOT IN operator
    - Example: Give all customers that have not placed any orders yet

    ```sql
    SELECT *
    FROM Customers
    WHERE CustomerID NOT IN (SELECT DISTINCT CustomerID FROM Orders)
    ```

    - This can also be accomplished with JOIN

    ```sql
    SELECT *
    FROM Customers c LEFT JOIN Orders o ON c.CustomerID = o.CustomerID
    WHERE o.CustomerID is NULL
    ```
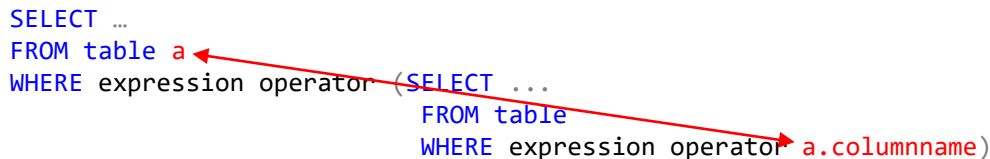
Correlated subqueries

- In a correlated subquery the inner query depends on information from the outer query.
  The subquery contains a search condition that refers to the main query, which makes the subquery depends on the main query

- The subquery is executed for each row in the main query. => $O(n^2)$

- The order of execution is from top to bottom, not from bottom to top as in a simple subquery, which is $O(n)$.

Correlated subqueries

- For performance reasons use joins or CTE (see further) or simple subquery if possible

- Principle

```
SELECT …
FROM table a
WHERE expression operator (SELECT ...
                           FROM table
                           WHERE expression operator a.columnname)
```

Correlated subqueries

- Example: Give employees with a salary larger than the average salary

```
SELECT FirstName + ' ' + LastName As FullName, Salary
FROM Employees
WHERE Salary > (SELECT AVG(Salary) FROM Employees)
```

This is a completely different (and more complex) task.

- Example: Give the employees whose salary is larger than the average of the salary of the employees <u>who report to the same boss</u>.

```
SELECT FirstName + ' ' + LastName As FullName, ReportsTo, Salary
FROM Employees As e
WHERE Salary > (SELECT AVG(Salary) FROM Employees WHERE ReportsTo = e.ReportsTo)
```

Remark: in the inner query you can use fields from the tables in the outer query but NOT vice versa.

0. Row 1 in the outer query

1. Outer query passes column values for that row to inner query

2. Inner query uses those values to evaluate inner query.

3. Inner query returns value to outer query, which decides if row in outer query will be kept.

4. This process repeats for each row in outer query.

Back to step 1.

HO GENT

Subqueries and the EXISTS operator

- The operator EXISTS tests the existence of a result set

  – Example: Give all customers that already placed an order

```
SELECT *
FROM Customers As c
WHERE EXISTS
(SELECT * FROM Orders WHERE CustomerID = c.customerID)
```

- There is also NOT EXISTS

  – Example: Give all customers that have not placed any orders yet

```
SELECT *
FROM Customers As c
WHERE NOT EXISTS
(SELECT * FROM Orders WHERE CustomerID = c.customerID)
```

HO
GENT

3 ways to accomplish the same result

- Example: Give all customers that did not place any orders yet
  - OUTER JOIN

```sql
SELECT *
FROM Customers c LEFT JOIN Orders o ON c.CustomerID = o.CustomerID
WHERE o.CustomerID is NULL
```

Which one will, in general, be the slowest?

  - Simple subquery

```sql
SELECT *
FROM Customers
WHERE CustomerID NOT IN (SELECT DISTINCT CustomerID FROM Orders)
```

  - Correlated subquery

```sql
SELECT *
FROM Customers As c
WHERE NOT EXISTS
(SELECT * FROM Orders WHERE CustomerID = c.customerID)
```

HO
GENT

Subqueries in the SELECT and FROM-clause

- The previous examples showed how subqueries can be used in the WHERE-clause

- Since the result of a query is a table it can be used in the FROM-clause.

- We will be using CTE's (see further) instead

HO
GENT

Subqueries in the SELECT-clause

- In a SELECT clause scalar (simple or correlated) subqueries can be used
  - Example: Give for each employee how much they earn more (or less) than the average salary of all employees with the same supervisor.
  - This example is just by way of example. We will be using CTE's instead.

```sql
SELECT Lastname, Firstname, Salary,
Salary -
(
    SELECT AVG(Salary)
    FROM Employees
    WHERE ReportsTo = e.ReportsTo
)
FROM Employees e
```

Some exercises

```
-- 1. Give the id and name of the products that have not been purchased yet.
-- 2. Select the names of the suppliers who supply products that have not been ordered yet.
-- 3. Give a list of all customers from the same country as the customer Maison Dewey
-- 4. Give for each product how much the price differs from the average price of all products of the same
category
-- 5. Give per title the employee that was last hired
-- 6. Which employee has processed most orders?
-- 7. What's the most common ContactTitle in Customers?
-- 8. Is there a supplier that has the same name as a customer?
-- 9. Give all the orders for which the ShipAddress is different from the CustomerAddress,
```

HO
GENT

DML

SQL - DML basic tasks

- SELECT → consulting data

- INSERT → adding data

- UPDATE → changing data

- DELETE → removing data

- MERGE → combine INSERT, UPDATE and DELETE

HO
GENT

Tip for not destroying your database

- The statements in this chapter are destructive.

- SQL has no UNDO by default!

- BUT you can 'simulate' UNDO if you take precautions.

Transactions are discussed in detail in one of the next chapters.

```sql
/* Tip for not destroying your database */
BEGIN TRANSACTION -- starts a new "transaction" -> Saves previous state of DB in buffer

-- several "destructive" commands can go here:
INSERT INTO Products(ProductName)
values ('TestProduct');

-- only you (in  your session) can see changes
SELECT * FROM Products WHERE ProductID = (SELECT MAX(ProductID) FROM Products)

ROLLBACK;   --> ends transaction and restores database in previous state
-- COMMIT;  --> ends transaction and makes changes permanent
```

HO
GENT

DML

INSERT: add new records

HO
GENT

Adding data - INSERT

- The INSERT statement adds data in a table
  - Add one row through via specification
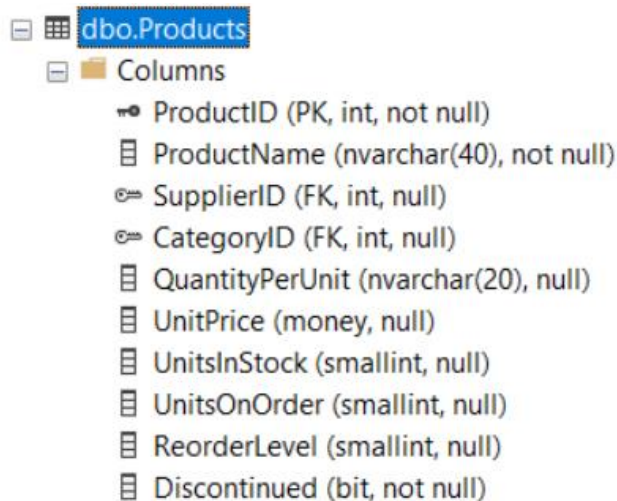 - Add selected row(s) from other tables

**HO
GENT**

INSERT of 1 row

- Method 1: specify only the (not NULL) values for specific columns.

- Method 2: specify all column values

- If the identity is generated automatically, this column can't be mentioned.

```sql
INSERT INTO Products (ProductName, CategoryID, Discontinued)
VALUES ('Toblerone', 3, 0)

INSERT INTO Products
VALUES ('Sultana', null, 3, null, null, null, null, null, 1)
```

HO
GENT

INSERT of 1 row

- The number of specified columns corresponds to the number of values.

- The specified values and corresponding columns have compatible data types.

- If no column names are specified the values are assigned in the column order as specified by the CREATE TABLE statement.

- Unmentioned columns get the value NULL or the DEFAULT value if any.

- NULL can also be specified as a value.

dbo.Products
  Columns
    ProductID (PK, int, not null)
    ProductName (nvarchar(40), not null)
    SupplierID (FK, int, null)
    CategoryID (FK, int, null)
    QuantityPerUnit (nvarchar(20), null)
    UnitPrice (money, null)
    UnitsInStock (smallint, null)
    UnitsOnOrder (smallint, null)
    ReorderLevel (smallint, null)
    Discontinued (bit, not null)

HO
GENT

INSERT of rows selected from other tables

- Mandatory fields have to be specified, unless they have a DEFAULT value.

- Constraints (see further) are validated.

- Unmentioned columns get the value NULL or the DEFAULT value if any.

```sql
INSERT INTO Customers (CustomerID, ContactName, ContactTitle, CompanyName)
SELECT substring(FirstName,1,2) + substring(LastName,1,3), FirstName + ' ' +
LastName, Title, 'EmployeeCompany'
FROM Employees
```

HO
GENT

DML

UPDATE: modify values

**HO
GENT**

Changing data - UPDATE

- Changing all rows in a table

  - Example: Increase the price of all products with 10%

```
UPDATE Products
SET UnitPrice = UnitPrice * 1.1
```

- Changing 1 row or a group of rows

  - Example: Increase the price of all
    the 'Bröd' products with 10%

```
UPDATE Products
SET UnitPrice = UnitPrice * 1.1
WHERE ProductName LIKE '%Bröd%'
```

  - Example: Increase the price of all the 'Bröd' products with 10%
    and set all units in stock to 0

```
UPDATE Products
SET UnitPrice = UnitPrice * 1.1, UnitsInStock = 0
WHERE ProductName LIKE '%Bröd%'
```

HO
GENT

Changing data - UPDATE

- Change rows based on data in another table
  - Standard SQL does not offer JOINs in an update statement → you can only use subqueries to refer to another table
  - Example: Due to a change in the euro – dollar exchange rate, we have to increase the unit price of products delivered by suppliers from the USA by 10%.

```
UPDATE Products
SET UnitPrice = (UnitPrice * 1.1)
WHERE SupplierID IN
(SELECT SupplierID FROM Suppliers WHERE Country = 'USA')
```

subquery

HO
GENT

DML

DELETE: remove records

HO GENT

Removing data - DELETE

- Delete some rows

  - Example: Delete the 'Bröd' products

    ```
    DELETE FROM Products
    WHERE ProductName LIKE '%Bröd%'
    ```

- Delete all rows in a table

  - via DELETE the identity values continues

    ```
    -- the identity value continues
    DELETE FROM Products
    ```

  - via TRUNCATE the identity value restarts from 1
    TRUNCATE is also more performant, but does not offer WHERE clause: it's all or nothing

    ```
    -- the identity value restarts from 1
    TRUNCATE TABLE Products
    ```

HO
GENT

DELETE - based on data in another table

- Change rows based on data in another table
  - Again no JOIN, only subquery
  - Example: Delete the orderdetails for all orders from the most recent orderdate

```
DELETE FROM OrderDetails
WHERE OrderID IN
(SELECT OrderID FROM Orders WHERE OrderDate = (SELECT MAX(OrderDate) from Orders))
```

HO
GENT

DML

MERGE: combine INSERT, UPDATE, DELETE

HO GENT

Merge

- With MERGE you can combine INSERT, UPDATE and DELETE.

- Very common use case: users work on an Excel sheet to update a relatively large amount of records because Excel offers a better overview than their ERP tool.

- They can update records, add new ones and delete records in Excel.

- After uploading the edited Excel file to a temporary table, the merge statement performs all UPDATEs, INSERTs and DELETEs at once.

HO
GENT

Merge

- Script to create temporary table

```
/* First execute following script to simulate the Excel file that has been imported
to a temporary table ShippersUpdate */

DROP TABLE IF EXISTS ShippersUpdate;
-- Add everything from Shippers to ShippersUpdate
SELECT * INTO ShippersUpdate FROM Shippers

-- Add an extra record to ShippersUpdate
INSERT INTO ShippersUpdate VALUES ('Pickup','(503) 555-9647')

-- Update a record of ShippersUpdate
UPDATE ShippersUpdate SET Phone = '(503) 555-4512' WHERE ShipperID = 1

-- Remove a record from ShippersUpdate
DELETE FROM ShippersUpdate WHERE shipperID = 4
```

HO
GENT

Merge

- Original table Shippers

Temporary table
             ShippersUpdate

	ShipperID	CompanyName	Phone
1	1	Speedy Express	(503) 555-9831
2	2	United Package	(503) 555-3199
3	3	Federal Shipping	(503) 555-9931
4	4	Total Shipping	(503) 555-9752
5	5	Federal Express	(503) 555-9773
6	7	PostNL	(503) 555-1236

	ShipperID	CompanyName	Phone
1	1	Speedy Express	(503) 555-4512
2	2	United Package	(503) 555-3199
3	3	Federal Shipping	(503) 555-9931
4	5	Federal Express	(503) 555-9773
5	6	Pickup	(503) 555-9647

Last row is added to Shippers at
the beginning of the script on the
next slide

Total Shipping is not in ShippersUpdate
Pickup is in ShippersUpdate and not in
Shippers

Merge

```
BEGIN TRANSACTION

INSERT INTO Shippers
VALUES ('PostNL', '(503) 555-1236')

SELECT * FROM Shippers;
SELECT * FROM ShippersUpdate;
```

HO
GENT

Merge

```sql
MERGE Shippers as t -- t = target
USING ShippersUpdate as s -- s = source
ON (t.ShipperID = s.ShipperID)

-- Which rows are in source and have different values for CompanyName or Phone?
-- Update those rows in target with the values coming from source
WHEN MATCHED AND t.CompanyName <> s.CompanyName OR ISNULL(t.Phone,'') <> ISNULL(s.Phone,'')
THEN UPDATE SET t.CompanyName = s.CompanyName, t.Phone=s.Phone

-- Which rows are in target and not in source?
-- Add those rows to source
WHEN NOT MATCHED BY target --> new rows
THEN INSERT (CompanyName, Phone) VALUES (s.CompanyName,s.Phone)

-- Which rows are in source and not in target?
-- Delete those rows from target
WHEN NOT MATCHED BY source  --> rows to delete
THEN DELETE;

-- Check the result
SELECT * FROM Shippers
ROLLBACK;
```

	ShipperID	CompanyName	Phone
1	1	Speedy Express	(503) 555-4512
2	2	United Package	(503) 555-3199
3	3	Federal Shipping	(503) 555-9931
4	5	Federal Express	(503) 555-9773
5	8	Pickup	(503) 555-9647

Remark: the option to delete records is a non standard extension by MS SQL.

Views

Views – Introduction

- Definition
    - A view is a saved SELECT statement
    - A view can be seen as a virtual table composed of other tables & views
    - No data is stored in the view itself, at each referral the underlying SELECT is re-executed;

HO GENT

Views – Introduction

- Advantages
  - Hide complexity of the database
    - Hide complex database design
    - Make large and complex queries accessible and reusable
    - Can be used as a partial solution for complex problems
  - Used for securing data access: revoke access to tables and grant access to customised views.
  - Organise data for export to other applications

HO GENT

Definition of a view

```
CREATE VIEW view_name [(column_list)]AS select_statement
```

- number of columns in (column_list) = # colunms in select
  - If no column names are specified, they are taken from the select
  - Column names are mandatory if the select statement contains calculations or joins in which some column names appear more than once
- the select statement may not contain an order by

Views – CRUD operations

- ## Creating a view

```sql
-- Creating a view
CREATE VIEW V_ProductsCustomer(productcode, company, quantity)
AS SELECT od.ProductID, c.CompanyName, sum(od.Quantity)
FROM Customers c
JOIN Orders o ON o.CustomerID = c.CustomerID
JOIN OrderDetails od ON o.OrderID = od.OrderID
GROUP BY od.ProductID, c.CompanyName;
```

- ## Using a view

```sql
-- Using a view
SELECT * FROM V_ProductsCustomer;
```

- ## Changing a view

```sql
-- Changing a view
ALTER VIEW V_ProductsCustomer(productcode, company)
AS SELECT od.ProductID, c.CompanyName
FROM Customers c
JOIN Orders o ON o.CustomerID = c.CustomerID
JOIN OrderDetails od ON o.OrderID = od.OrderID
GROUP BY od.ProductID, c.CompanyName;
```

- ## Deleting a view

```sql
-- Dropping a view
DROP VIEW V_ProductsCustomer;
```

Example: views as partial solution for complex problems

- Example: Create a view with the number of orders per employee per year

- Example: Calculate per employee and per year the running total of processed orders.

```
-- Create a view with the number of orders per employee per year
CREATE OR ALTER VIEW vw_number_of_orders_per_employee_per_year
AS
SELECT EmployeeID, YEAR(OrderDate) As OrderYear, COUNT(OrderID) As NumberOfOrders
FROM Orders
GROUP BY EmployeeID, YEAR(OrderDate)

-- Check the result
SELECT * FROM vw_number_of_orders_per_employee_per_year
```

Example: views as partial solution for complex problems

```
-- Add the running total. Use the view.
SELECT EmployeeID, OrderYear, NumberOfOrders,
(SELECT SUM(NumberOfOrders)
FROM vw_number_of_orders_per_employee_per_year
WHERE OrderYear <= vw.OrderYear AND EmployeeID = vw.EmployeeID
) As TotalNumberOfOrders
FROM vw_number_of_orders_per_employee_per_year vw
ORDER BY EmployeeID, OrderYear
```

- Drawback of using views in this way:
  views are stored in the database and might create a mess if
  you have hundreds of them.
  We'll try to solve this using cte's.

Views in SQL Server Management Studio

- Simple views can also be made with the graphical user interface

- This is not possible for views with subqueries or common table expressions


```
SELECT ProductID, ProductName, SupplierID, CategoryID, QuantityPerUnit,
UnitPrice, UnitsInStock, UnitsOnOrder, ReorderLevel, Discontinued
FROM    dbo.Products
WHERE  (CategoryID = 1)
```

Views in SQL Server Management Studio

- The easiest way to change the SQL code of an existing view is by right clicking on the name of the view and:
  - Script View as…
  - ALTER to…
  - New Query Editor Window

Exercises

```
-- Exercise 1
-- The company wants to weekly check the stock of their products.
-- If the stock is below 15, they'd like to order more to fulfill the need.

-- (1.1) Create a QUERY that shows the ProductId, ProductName and the name of the supplier, do not forget the
WHERE clause.
-- (1.2) Turn this SELECT statement into a VIEW called: vw_products_to_order.
-- (1.3) Query the VIEW to see the results.

-- Exercise 2
-- The company has to increase prices of certain products. To make it seem the prices are not increasing
dramatically they're planning to spread the price increase over multiple years. In total they'd like a 10%
price for certain products. The list of impacted products can grow over the coming years.
-- We'd like to keep all the logic of selecting the correct products in 1 SQL View, in programming terms
'keeping it DRY'.
-- The updating of the items is not part of the view itself.
-- The products in scope are all the products with the term 'Bröd' or 'Biscuit'.

-- (2.1) Create a simple SQL Query to get the correct resultset
-- (2.2) Turn this SELECT statement into a VIEW called: vw_price_increasing_products.
-- (2.3) Query the VIEW to see the results.
```

COMMON TABLE EXPRESSIONS

HO
GENT

Common Table Expressions:
the WITH component

- Example: Give per category the minimum price and all products with that minimum price

```
-- Solution 1 -> with subqueries

SELECT CategoryID, ProductID, UnitPrice
FROM Products p
WHERE UnitPrice = (SELECT MIN(UnitPrice) FROM Products WHERE CategoryID = p.CategoryID)
```

- Not performant! Loops through all products and calculates the MIN(unitprice) for the category of that specific product: O(n²)

- The MIN(unitprice) is calculated multiple times for each category!

Common Table Expressions: the WITH component

- Example: Give per category the minimum price and all products with that minimum price

- Solution 2 → CTE's (**C**ommon **T**able **E**xpression)

```
-- Solution 2 -> with CTE's
WITH CategoryMinPrice(CategoryID, MinPrice)
AS (SELECT CategoryID, MIN(UnitPrice)
    FROM Products AS p
    GROUP BY CategoryID)

SELECT c.CategoryID, p.ProductID, MinPrice
FROM Products AS p
JOIN CategoryMinPrice AS c ON p.CategoryID = c.CategoryID AND p.UnitPrice = c.MinPrice;
```

Common Table Expressions: the WITH component

- Solution 2 → CTE's (**C**ommon **T**able **E**xpression)

```sql
-- Solution 2 -> with CTE's
WITH CategoryMinPrice(CategoryID, MinPrice)
AS (SELECT CategoryID, MIN(UnitPrice)
    FROM Products AS p
    GROUP BY CategoryID)

SELECT c.CategoryID, p.ProductID, MinPrice
FROM Products AS p
JOIN CategoryMinPrice AS c ON p.CategoryID = c.CategoryID AND p.UnitPrice = c.MinPrice;
```

- Using the WITH-component you can give the subquery its own name (with column names) and reuse it in the rest of the query (possibly several times!)

Common Table Expressions: the WITH component

The columns in the CTE should have a name, so
you can refer to these columns.
(1) If not given a name, it will use the 'default' name (e.g. CategoryID, MIN(UnitPrice))
(2) Or you can specify the columnname in the 'header' of the CTE
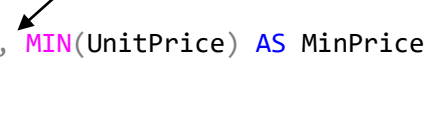(3) Or you can give each column a name in the CTE.

```sql
-- Solution 2
WITH CategoryMinPrice(CategoryID, MinPrice)
AS (SELECT CategoryID, MIN(UnitPrice)
    FROM Products AS p
    GROUP BY CategoryID)

SELECT c.CategoryID, p.ProductID, MinPrice
FROM Products AS p
JOIN CategoryMinPrice AS c ON p.CategoryID = c.CategoryID AND p.UnitPrice = c.MinPrice;
```

```sql
-- Solution 2
WITH CategoryMinPrice
AS (SELECT CategoryID As CategoryID, MIN(UnitPrice) AS MinPrice
    FROM Products AS p
    GROUP BY CategoryID)

SELECT c.CategoryID, p.ProductID, MinPrice
FROM Products AS p
JOIN CategoryMinPrice AS c ON p.CategoryID = c.CategoryID AND p.UnitPrice = c.MinPrice;
```

Common Table Expressions: the WITH component

- the WITH-component has two application areas:
  - Simplify SQL-instructions, e.g. simplified alternative for simple subqueries or avoid repetition of SQL constructs
  - Traverse recursively hierarchical and network structures

CTE's versus Views

- Similarities
  - WITH ~ CREATE VIEW
  - Both are virtual tables:
    the content is derived from other tables
- Differences
  - A CTE only exists during the SELECT-statement
  - A CTE is not visible for other users and applications

CTE's versus Subqueries

- Similarities
  - Both are virtual tables:
    the content is derived from other tables

- Differences
  - A CTE can be reused in the same query
  - A subquery is defined in the clause where it is used (SELECT/FROM/WHERE/…)
  - A CTE is defined on top of the query
  - A simple subquery can always be replaced by a CTE

CTE's with more than 1 WITH - component

- Example: Give per year per customer the relative contribution of this customer to the total revenue

- Step 1: Calculate the total revenue per year

```sql
-- Step 1 -> Total revenue per year

SELECT YEAR(OrderDate), SUM(od.UnitPrice * od.Quantity)
FROM Orders o INNER JOIN OrderDetails od
ON o.OrderID = od.OrderID
GROUP BY YEAR(OrderDate)
```

CTE's with more than 1 WITH - component

- Step 2: Calculate the total revenue per year per customer

```
-- Step 2 -> Total revenue per year per customer

SELECT YEAR(OrderDate), o.CustomerID, SUM(od.UnitPrice * od.Quantity)
FROM Orders o INNER JOIN OrderDetails od
ON o.OrderID = od.OrderID
GROUP BY YEAR(OrderDate), o.CustomerID
```

CTE's with more than 1 WITH - component

- Step 3: Combine both

```sql
-- Step 3 -> Combine both
WITH TotalRevenuePerYear(RevenueYear, TotalRevenue)
AS
(SELECT YEAR(OrderDate), SUM(od.UnitPrice * od.Quantity)
FROM Orders o INNER JOIN OrderDetails od
ON o.OrderID = od.OrderID
GROUP BY YEAR(OrderDate)),

TotalRevenuePerYearPerCustomer(RevenueYear, CustomerID, Revenue)
AS
(SELECT YEAR(OrderDate), o.CustomerID, SUM(od.UnitPrice * od.Quantity)
FROM Orders o INNER JOIN OrderDetails od
ON o.OrderID = od.OrderID
GROUP BY YEAR(OrderDate), o.CustomerID)

SELECT CustomerID, pc.RevenueYear, FORMAT(Revenue / TotalRevenue, 'P') As Relative
FROM TotalRevenuePerYearPerCustomer pc INNER JOIN TotalRevenuePerYear t
ON pc.RevenueYear = t.RevenueYear
ORDER BY 2 ASC, 3 DESC
```

	CustomerID	RevenueYear	RelativePart
1	ERNSH	2016	7.58%
2	QUICK	2016	5.62%
3	QUEEN	2016	5.42%
4	PICCO	2016	5.38%
5	SAVEA	2016	5.38%
6	RATTC	2016	4.79%
7	FRANK	2016	4.72%
8	HUNGO	2016	4.66%
9	BLONP	2016	4.41%
10	SPLIR	2016	3.71%
11	SUPRD	2016	2.84%
12	MEREP	2016	2.70%
13	SEVES	2016	2.62%
14	LILAS	2016	2.59%
15	OLDWO	2016	2.24%

Recursive SELECT's

- Recursivity means:
  - We continue to execute a table expression <u>until a condition</u> is reached.
- This allows you to solve problems like:
  - Who are the friends of my friends etc. (in a social network)?
  - What is the hierarchy of an organisation ?
  - Find the parts and subparts of a product (Bill of materials).

Recursive SELECT's

- Example:
  Give the integers from 1 to 5

```sql
WITH numbers(number) AS
(SELECT 1
  UNION all
   SELECT number + 1
   FROM numbers
   WHERE number < 5)

SELECT * FROM numbers;
```

- Characteristics of recursive use of WITH:
  - The with component consists of (at least) 2 expressions, combined with union all
  - A temporary table is consulted in the second expression
  - At least one of the expressions may not refer to the temporary table.

Recursive SELECT's: how does it work?

- 1. SQL searches the table expressions that don't contain recursivity and executes them one by one.

```
SELECT 1
```

	number
1	1

- 2. Execute all recursive expressoins. The numbers table, that got a value of 1 in step 1, is used.
  This row is added to the numbers table.

```
SELECT number + 1
FROM numbers
WHERE number < 5
```

	number
1	2

Recursive SELECT's: how does it work?

- 3. Now the recursion starts: the 2nd expression is re-executed, giving as result:
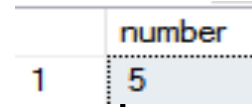Remark: not all rows added in the previous steps are processed, but only those rows (1 row in this example), that were added in the previous step (step 2).

- 4. Since step 3 also gave a result, the recursive expression is executed again, producing as intermediate result:

Recursive SELECT's: how does it work?

- 5. And this happens again:

	number
1	5

- If the expression is now processed again, it does not return a result, since in the previous step no rows were added that correspond to the condition number < 5.
Here SQL stops the processing of the table expression and the final result is known.

- Summary: the 1st (non-recursive) expression is executed once and the 2nd expression is executed until it does not return any more results.

Recursive SELECT's: max number of recursions = 100

- Example: Give the numbers from 1 to 999

```sql
WITH numbers(number) AS
(SELECT 1
  UNION all
    SELECT number + 1
    FROM numbers
    WHERE number < 999)

SELECT * FROM numbers;
```

Msg 530, Level 16, State 1, Line 260
The statement terminated. The maximum recursion 100 has been exhausted before statement completion.

Recursive SELECT's: OPTION maxrecursion

- Example: Give the numbers from 1 to 999

```
WITH numbers(number) AS
(SELECT 1
  UNION all
    SELECT number + 1
    FROM numbers
    WHERE number < 999)

SELECT * FROM numbers
OPTION (maxrecursion 1000);
```

Application: generate missing months

- Example: Give the total revenue per month in 2016
  Not all months occur

```
SELECT YEAR(OrderDate)*100 + Month(OrderDate) AS RevenueMonth, SUM(od.UnitPrice * od.Quantity) AS Revenue
FROM Orders o INNER JOIN OrderDetails od ON o.OrderID = od.OrderID
WHERE YEAR(OrderDate) = 2016
GROUP BY YEAR(OrderDate)*100 + Month(OrderDate)
```

	RevenueMonth	Revenue
1	201607	30192,10
2	201608	26609,40
3	201609	27636,00
4	201610	41203,60
5	201611	49704,00
6	201612	50953,40

Application: generate missing months

- Solution: Generate all months with CTE

```
WITH Months AS
(SELECT 201601 as RevenueMonth
UNION ALL
SELECT RevenueMonth + 1
FROM Months
WHERE RevenueMonth < 201612)
SELECT * FROM Months;
```

	RevenueMonth
1	201601
2	201602
3	201603
4	201604
5	201605
6	201606
7	201607
8	201608
9	201609
10	201610
11	201611
12	201612

Application: generate missing months

- Solution: … and combine with LEFT JOIN

```sql
WITH Months(RevenueMonth) AS
(SELECT 201601 as RevenueMonth
UNION ALL
SELECT RevenueMonth + 1
FROM Months
WHERE RevenueMonth < 201612),

Revenues(RevenueMonth, Revenue)
AS
(SELECT YEAR(OrderDate)*100 + Month(OrderDate) AS RevenueMonth,
SUM(od.UnitPrice * od.Quantity) AS Revenue
FROM Orders o INNER JOIN OrderDetails od ON o.OrderID = od.OrderID
WHERE YEAR(OrderDate) = 2016
GROUP BY YEAR(OrderDate)*100 + Month(OrderDate))

SELECT m.RevenueMonth, ISNULL(r.Revenue, 0) As Revenue
FROM Months m LEFT JOIN Revenues r ON m.RevenueMonth = r.RevenueMonth
```

	RevenueMonth	Revenue
1	201601	0,00
2	201602	0,00
3	201603	0,00
4	201604	0,00
5	201605	0,00
6	201606	0,00
7	201607	30192,10
8	201608	26609,40
9	201609	27636,00
10	201610	41203,60
11	201611	49704,00
12	201612	50953,40

Recursively traversing a hierarchical structure

- Example: Give all employees who report directly or indirectly to Andrew Fuller (ReportsTo IS NULL)
  - Step 1 returns all employees that report directly to Andrew Fuller
  - Step 2 adds the second 'layer':
    who reports to
    someone who
    reports to
    A. Fuller

```
WITH Bosses (boss, emp)
AS
(SELECT ReportsTo, EmployeeID
FROM Employees
WHERE ReportsTo IS NULL
UNION ALL
SELECT e.ReportsTo, e.EmployeeID
FROM Employees e INNER JOIN Bosses b ON e.ReportsTo = b.emp)

SELECT * FROM Bosses
ORDER BY boss, emp;
```

	boss	emp
1	NULL	2
2	2	1
3	2	3
4	2	4
5	2	5
6	2	8
7	5	6
8	5	7
9	5	9

Recursively traversing a hierarchical structure

- Draw the organization chart.

	boss	emp	title	level	path
1	NULL	2	Vice President, Sales	1	Vice President, Sales
2	2	1	Sales Representative	2	Vice President, Sales<--Sales Representative
3	2	3	Sales Representative	2	Vice President, Sales<--Sales Representative
4	2	4	Sales Representative	2	Vice President, Sales<--Sales Representative
5	2	5	Sales Manager	2	Vice President, Sales<--Sales Manager
6	2	8	Inside Sales Coordinator	2	Vice President, Sales<--Inside Sales Coordinator
7	5	6	Sales Representative	3	Vice President, Sales<--Sales Manager<--Sales Representative
8	5	7	Sales Representative	3	Vice President, Sales<--Sales Manager<--Sales Representative
9	5	9	Sales Representative	3	Vice President, Sales<--Sales Manager<--Sales Representative

Recursively traversing a hierarchical structure

- Draw the organization chart.

```
WITH Bosses (boss, emp, title, level, path)
AS
(SELECT ReportsTo, EmployeeID, Title, 1, convert(varchar(max), Title)
FROM Employees
WHERE ReportsTo IS NULL
UNION ALL
SELECT e.ReportsTo, e.EmployeeID, e.Title, b.level + 1, convert(varchar(max), b.path + '<--' + e.title)
FROM Employees e INNER JOIN Bosses b ON e.ReportsTo = b.emp)

SELECT * FROM Bosses
ORDER BY boss, emp;
```

Exercises

```
-- 1. Give all employees that started working as an employee in the same year as Robert King

-- 2 Make a histogram of the number of orders per customer, so show how many times each number occurs.
-- E.g. in the graph below: 1 customer placed 1 order, 2 customers placed 2 orders, 7 customers placed 3
orders, etc.

/*

nrNumberOfCustomers
11
22
37
46
510
68
77
...

*/
```

Exercises

```
-- 3. Give the customers of the Country in which most customers live
-- 4. Give all employees except for the eldest. Solve this first using a subquery and afterwards using a cte
-- 5. What is the total number of customers and suppliers?
-- 6. Give per title the eldest employee
-- 7. Give per title the employee that earns most
-- 8. Give the titles for which the eldest employee is also the employee who earns most
```

Exercises

```
-- 9. Execute the following script:
CREATE TABLE Parts
(
    [Super]   CHAR(3) NOT NULL,
    [Sub]     CHAR(3) NOT NULL,
    [Amount]  INT NOT NULL,
    PRIMARY KEY(Super, Sub)
);

INSERT INTO Parts VALUES ('O1','O2',10);
INSERT INTO Parts VALUES ('O1','O3',5);
INSERT INTO Parts VALUES ('O1','O4',10);
INSERT INTO Parts VALUES ('O2','O5',25);
INSERT INTO Parts VALUES ('O2','O6',5);
INSERT INTO Parts VALUES ('O3','O7',10);
INSERT INTO Parts VALUES ('O6','O8',15);
INSERT INTO Parts VALUES ('O8','O11',5);
INSERT INTO Parts VALUES ('O9','O10',20);
INSERT INTO Parts VALUES ('O10','O11',25);

-- Show all parts that are directly or indirectly part of O2,
-- so all parts of which O2 is composed.
-- Add an extra column with the path as shown:
```

```
/*
SUPER  SUB    PAD
O2     O5     O2 <-O5
O2     O6     O2 <-O6
O6     O8     O2 <-O6 <-O8
O8     O11    O2 <-O6 <-O8 <-O11
*/
```